

Open Source Application Performance Monitoring (APM)

*Ein Überblick über APM Tools und Standards für
Java-basierte Enterprise-Anwendungen*

Dr. Andreas Brunnert,
Bernhard Lubomski
RETIT GmbH



Motivation

The amount of open source APM tools has grown dramatically in the last four years:



PinPoint



Jaeger



Zipkin



inspectIT



elasticAPM



**Apache
Skywalking**



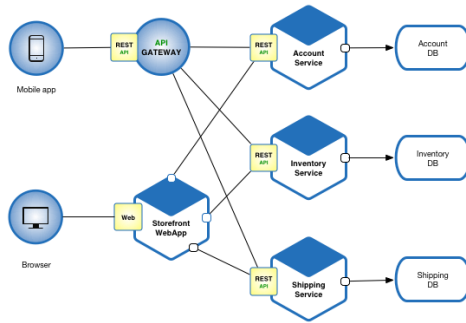
**Stage-
monitor**



**App-
dash**

Motivation

Complexity increase in modern software systems



Services might need to interact with each other in ways that might not be obvious at the time of development or deployment.

Growing importance of IT for more business models



Downtimes or bad software performance have a direct impact on revenue.

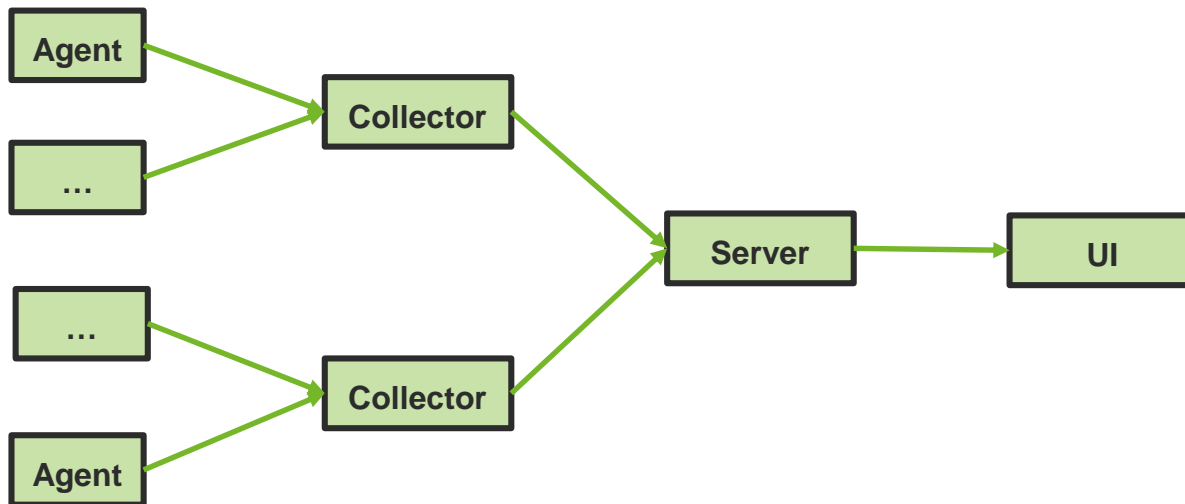
Development of tracing standards



Which allow to easily exchange the tracing tool in use. Furthermore, they reduce the effort for each vendor.

Context

Anatomy of an APM Solution

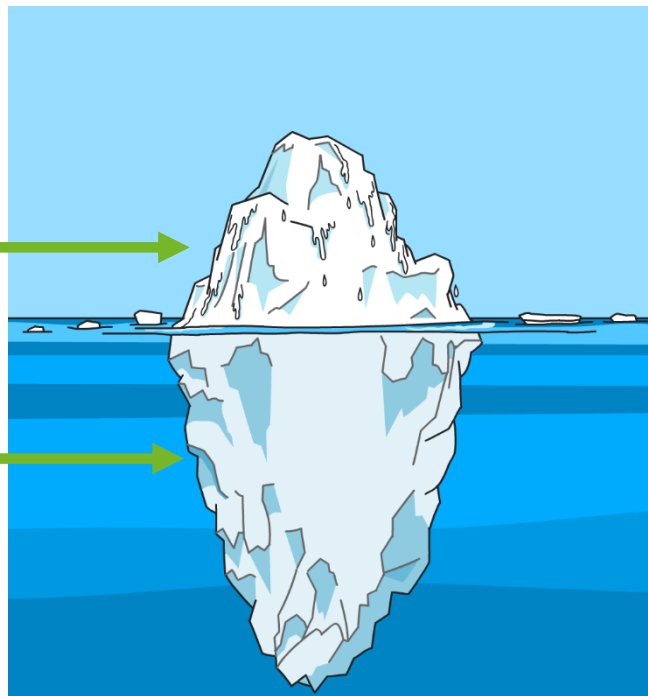


Context

Code and Effort distribution of an APM Solution

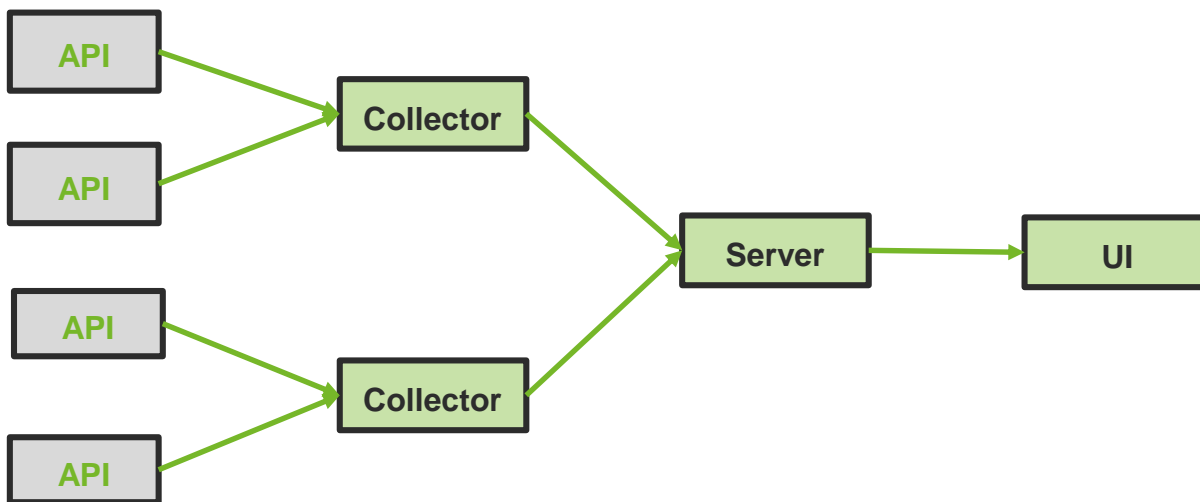
UI + Server + Collectors

Agents



Context

Scope of many open source APM solutions



RESEARCH › PUBLICATIONS ›

Dapper, a Large-Scale Distributed Systems Tracing Infrastructure



Download



Search



Copy Bibtex

- Some tools build upon the same concepts or even fork each other:
 - <https://research.google.com/pubs/pub36356.html>
 - Basis for: Pinpoint, Jaeger and Zipkin
 - Zipkin is again the basis for Jaeger

Context

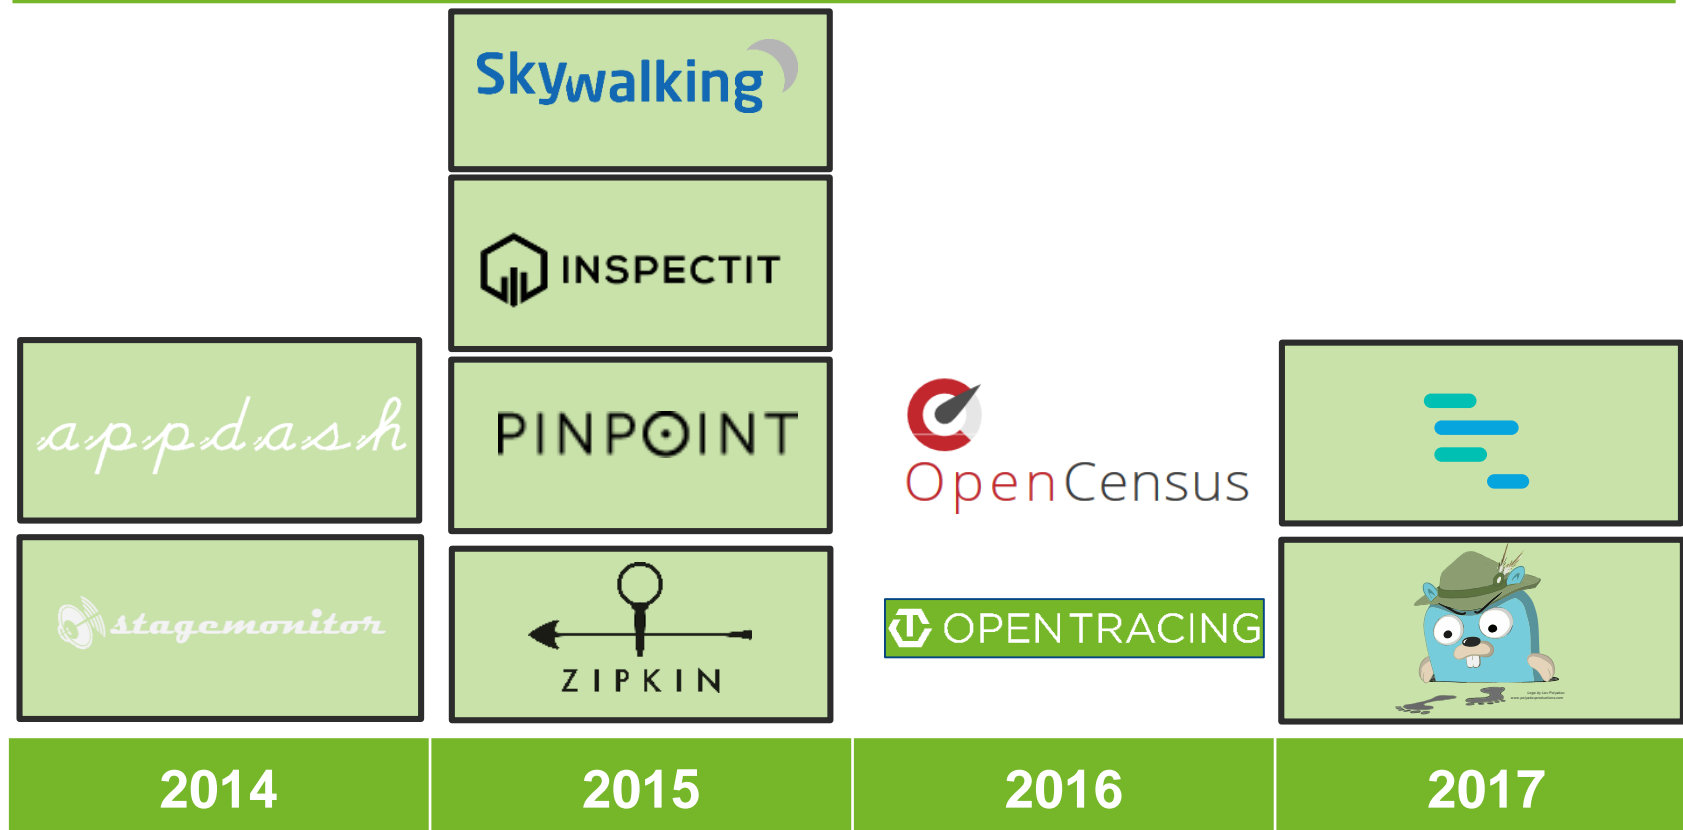
But how do these open source APM tools compare?

- Age
- Popularity
- Supported Technologies
- Standards Support

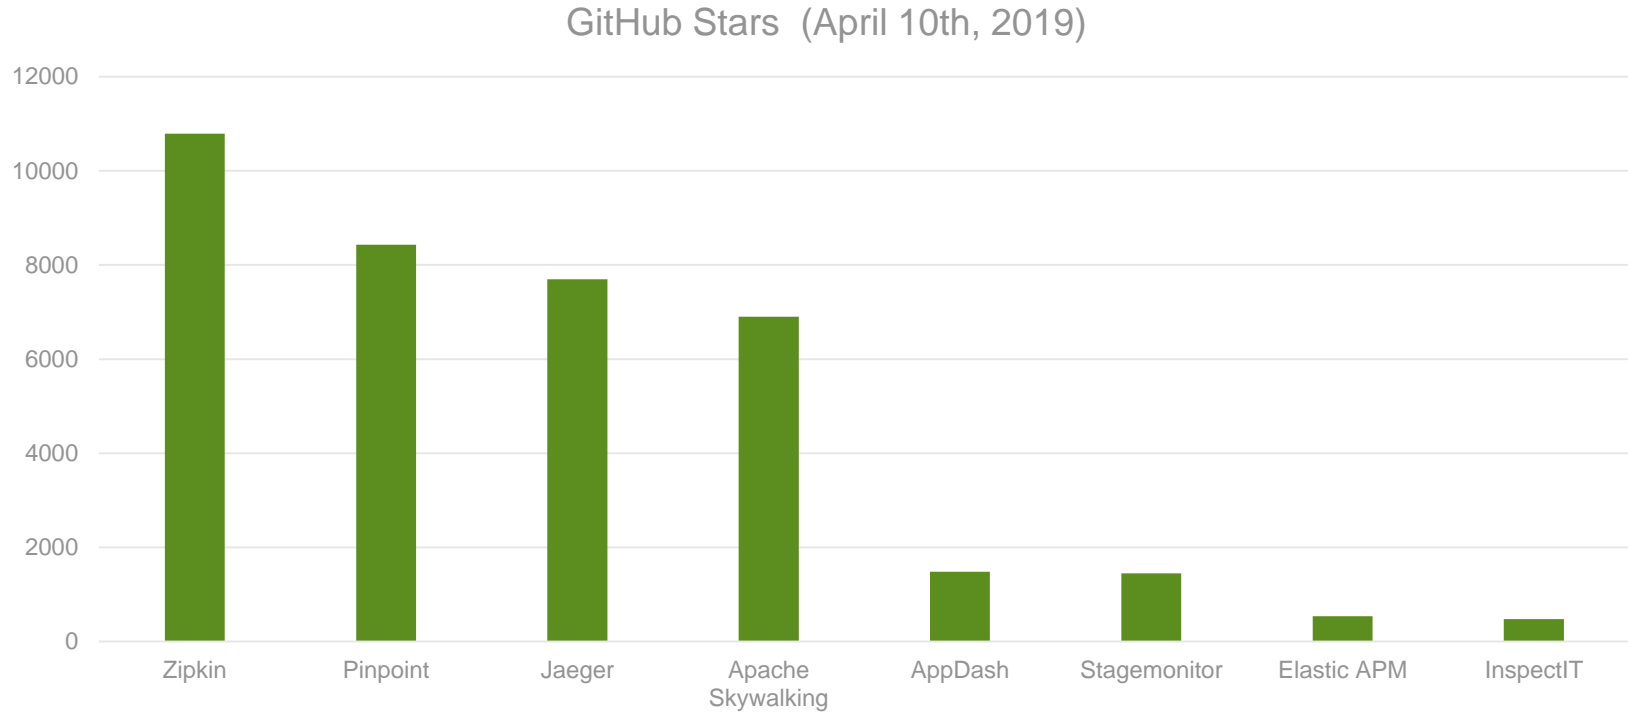
- Not in presentation (will be covered in later blog articles):
 - Setup – Effort
 - Integration Capabilities with other tools
 - License

What are reasons for a closed source alternative?

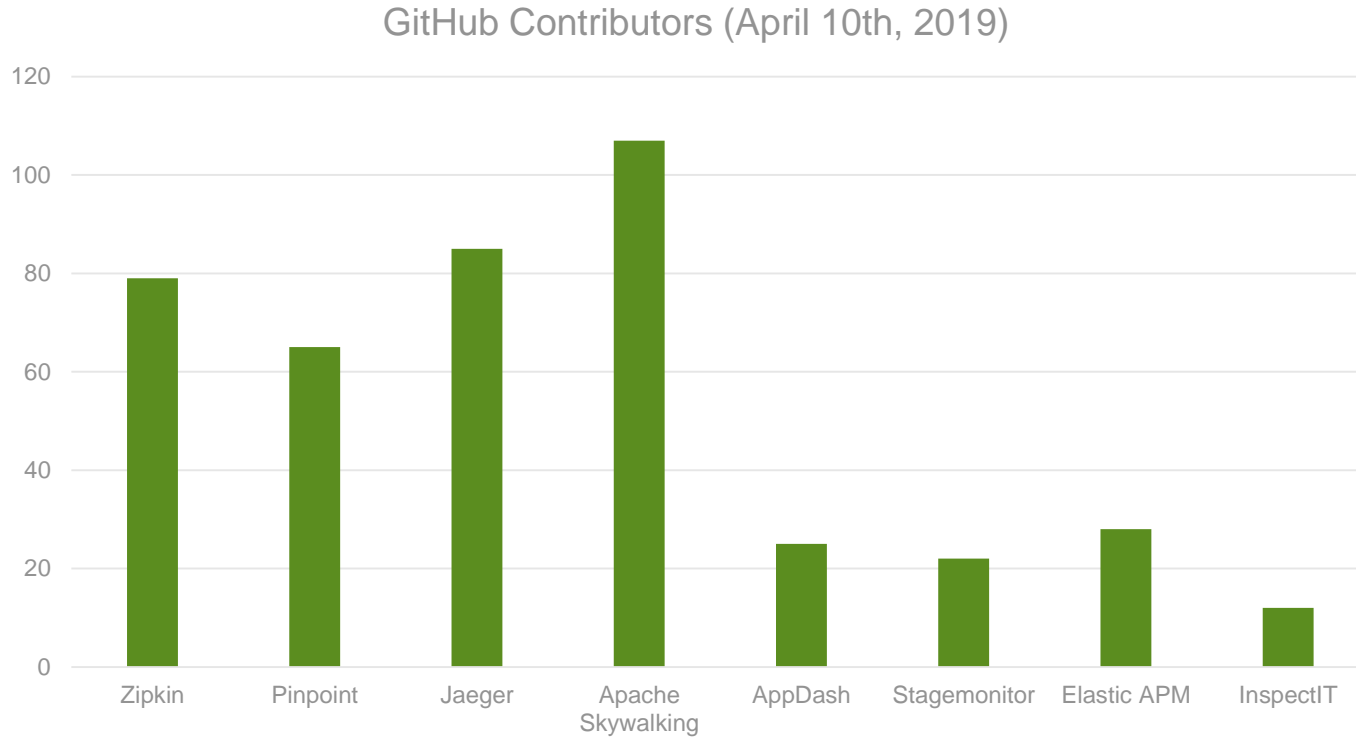
A brief timeline of tool availability (since 2014)



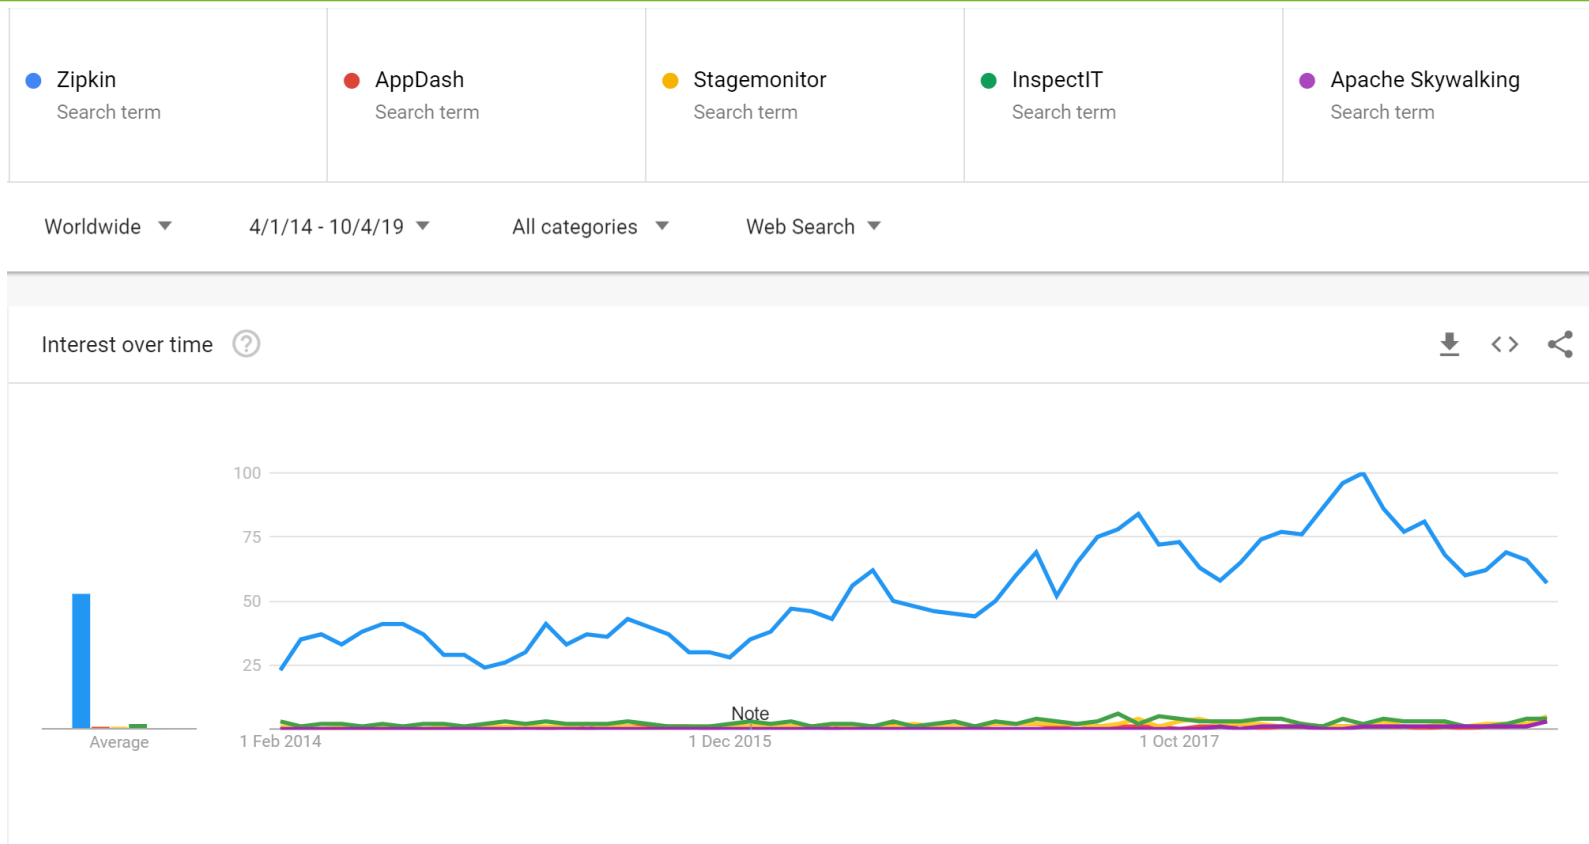
A ranking of GitHub stars



A ranking of GitHub contributes



Google Trends Analysis

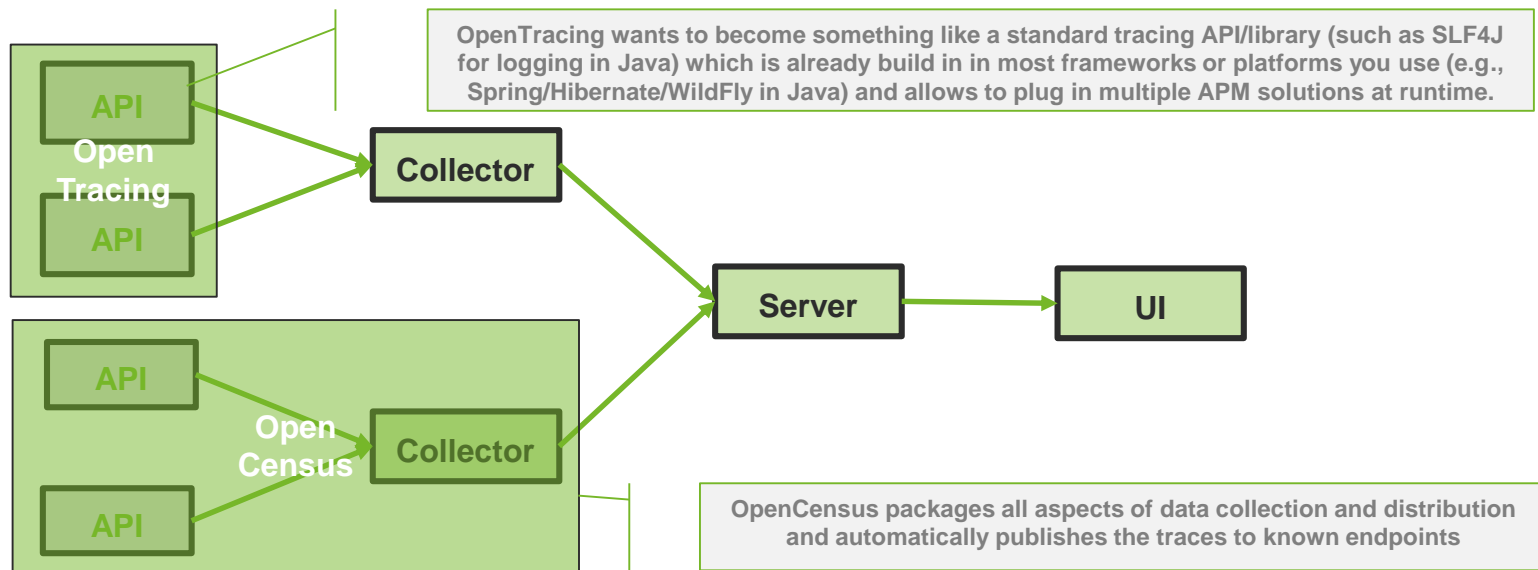


Open Source “Standards”

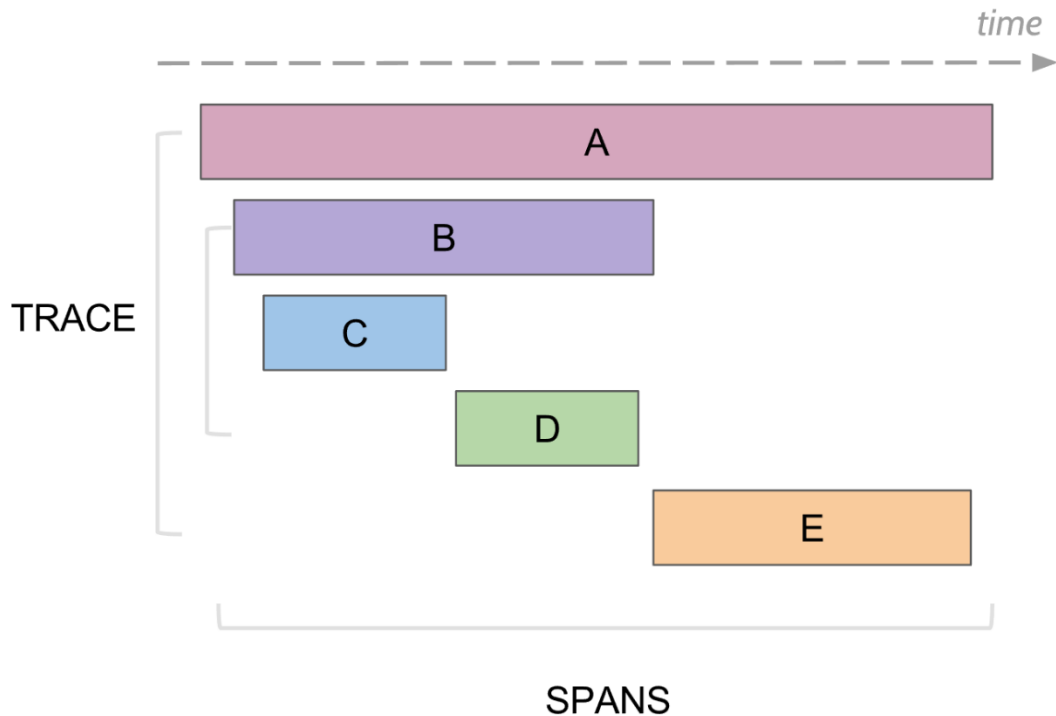


Open Source “Standards”

Scope of OpenTracing vs. OpenCensus (Simplified)



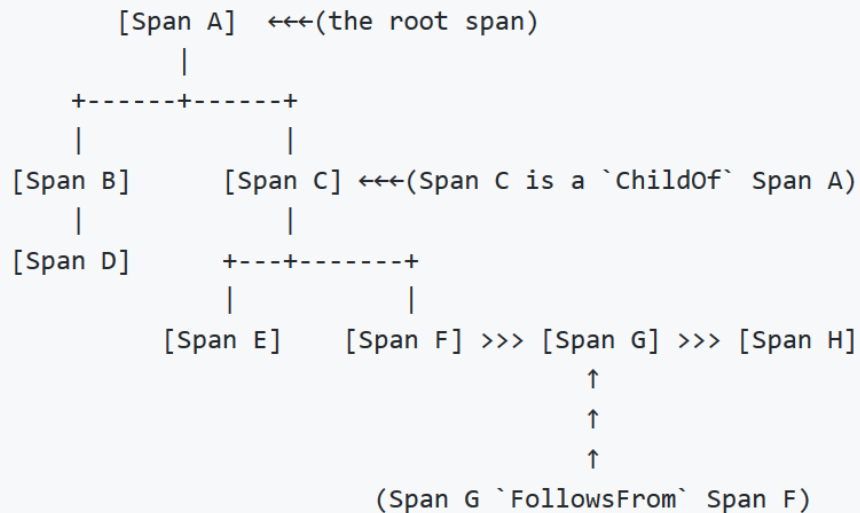
Open Source “Standards” - OpenTracing



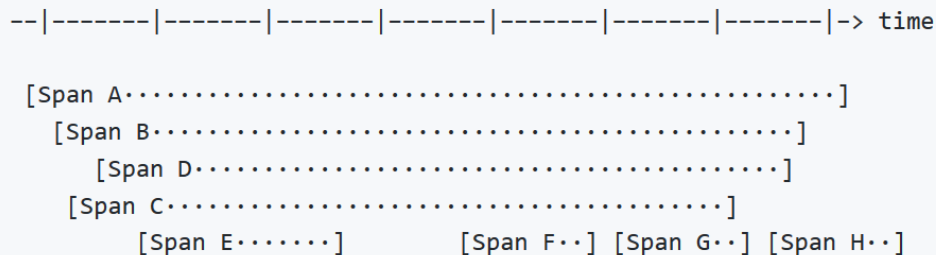
Source: <https://www.jaegertracing.io/docs/architecture/>

Open Source “Standards” - OpenTracing

Causal relationships between Spans in a single Trace



Temporal relationships between Spans in a single Trace



Source: <https://github.com/opentracing/specification/blob/master/specification.md>

Open Source “Standards” - OpenTracing

```
import io.jaegertracing.Configuration;
import io.opentracing.Span;
import io.opentracing.util.GlobalTracer;

...
```

You only need to do this once

```
GlobalTracer.register(
    Configuration.fromEnv().getTracer());
);
```

Tracer configuration loaded from environment properties, but can be customized programmatically

```
...

try (Scope scope = tracer.buildSpan("parentSpan").startActive()) {
    try (Scope scope = tracer.buildSpan("childSpan").startActive()) {
        // "child" is automatically a child of "parent".
    }
}
```

For each individual span

Source:

<https://github.com/jaegertracing/jaeger-client-java/blob/master/jaeger-core/README.md>

<https://opentracing.io/guides/java/>

Open Source “Standards” - OpenCensus

Languages

LANGUAGE	STATS	TRACING
C#	Supported	Supported
C++	Supported	Supported
Erlang/Elixir	Supported	Supported
Go	Supported	Supported
Java (JVM, OpenJDK)	Supported	Supported
Node.js	Supported	Supported
PHP	Planned	Supported
Python	Supported	Supported
Ruby	Planned	Supported

Source:

<https://opencensus.io/roadmap/index.html>

Open Source “Standards” - OpenCensus

Exporters

T Backend supports Tracing

S Backend supports Stats

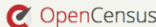
BACKEND	C#	C++	ERLANG	GO	JAVA	NODE.JS	PHP	PYTHON	RUBY
AWS X-Ray	-	-	-	T		-	-	-	-
Azure Monitor	T S	-	-	T	T	-	-	T	-
Datadog	-	-	T S	T S	T	-	-	-	-
Elasticsearch	-	-	-	-	T	-	-	-	-
Honeycomb	-	-	-	T	-	-	-	-	-
Instana	-	-	-	-	T	T	-	-	-
Jaeger	-	-	-	T	T	T	-	T	-
Prometheus	S	S	S	S	S	S	-	S	-
SignalFX	-	-	-	-	S	-	-	-	-
Stackdriver	T	T S	T	T S	T S	T S	-	T S	-
Zipkin	T	T	T	T	T	T	-	T	-

Source:

<https://opencensus.io/roadmap/index.html>

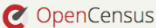
Open Source “Standards” - OpenCensus

- zPages: in process web pages, displaying collected data from process
- No backend necessary.
- Useful for debugging.
- Available for Go, Java and Node.js.



TraceZ Summary

Span Name	Running	Latency Samples										Error Samples
		[>0us]	[>10us]	[>100us]	[>1ms]	[>10ms]	[>100ms]	[>1s]	[>10s]	[>100s]		
ExampleSpan	0	0	0	0	0	0	1	0	0	0	0	
HttpServer/rpcz	0	0	0	0	0	1	0	0	0	0	0	
HttpServer/statusz	0	0	0	0	2	3	0	0	0	0	0	
HttpServer/traceconfigz	0	0	0	0	0	1	0	0	0	0	0	
HttpServer/tracez	0	0	0	0	0	0	0	0	0	0	0	



RPC Stats

Sent

Method	Count			Avg latency (ms)			Rate (rpc/s)			Input (kb/s)			Output (kb/s)			Errors		
	Min.	Hr.	Tot.	Min.	Hr.	Tot.	Min.	Hr.	Tot.	Min.	Hr.	Tot.	Min.	Hr.	Tot.	Min.	Hr.	Tot.
ExampleMethod/abc.123/xyz.890/example	1	1	1	1.000	1.000	1.000	0.017	0.000	0.042	1.628	0.027	4.123	1.628	0.027	4.122	1	1	1

Received

Method	Count			Avg latency (ms)			Rate (rpc/s)			Input (kb/s)			Output (kb/s)			Errors		
	Min.	Hr.	Tot.	Min.	Hr.	Tot.	Min.	Hr.	Tot.	Min.	Hr.	Tot.	Min.	Hr.	Tot.	Min.	Hr.	Tot.
ExampleMethod/abc.123/xyz.890/example	1	1	1	1.000	1.000	1.000	0.017	0.000	0.042	1.628	0.027	4.118	1.628	0.027	4.118	1	1	1

Source: <https://opencensus.io/zpages/#zpages>

Open Source “Standards” - OpenCensus

```
import io.opencensus.common.Scop
import io.opencensus.exporter.trace.zipkin.ZipkinTraceExporter;
import io.opencensus.trace.Tracer;
```

You only need to do this once

```
...
ZipkinTraceExporter.createAndRegister("http://127.0.0.1:9411/api/v2/spans", "my-service");
Tracer tracer = Tracing.getTracer(); // Global singleton Tracer object
```

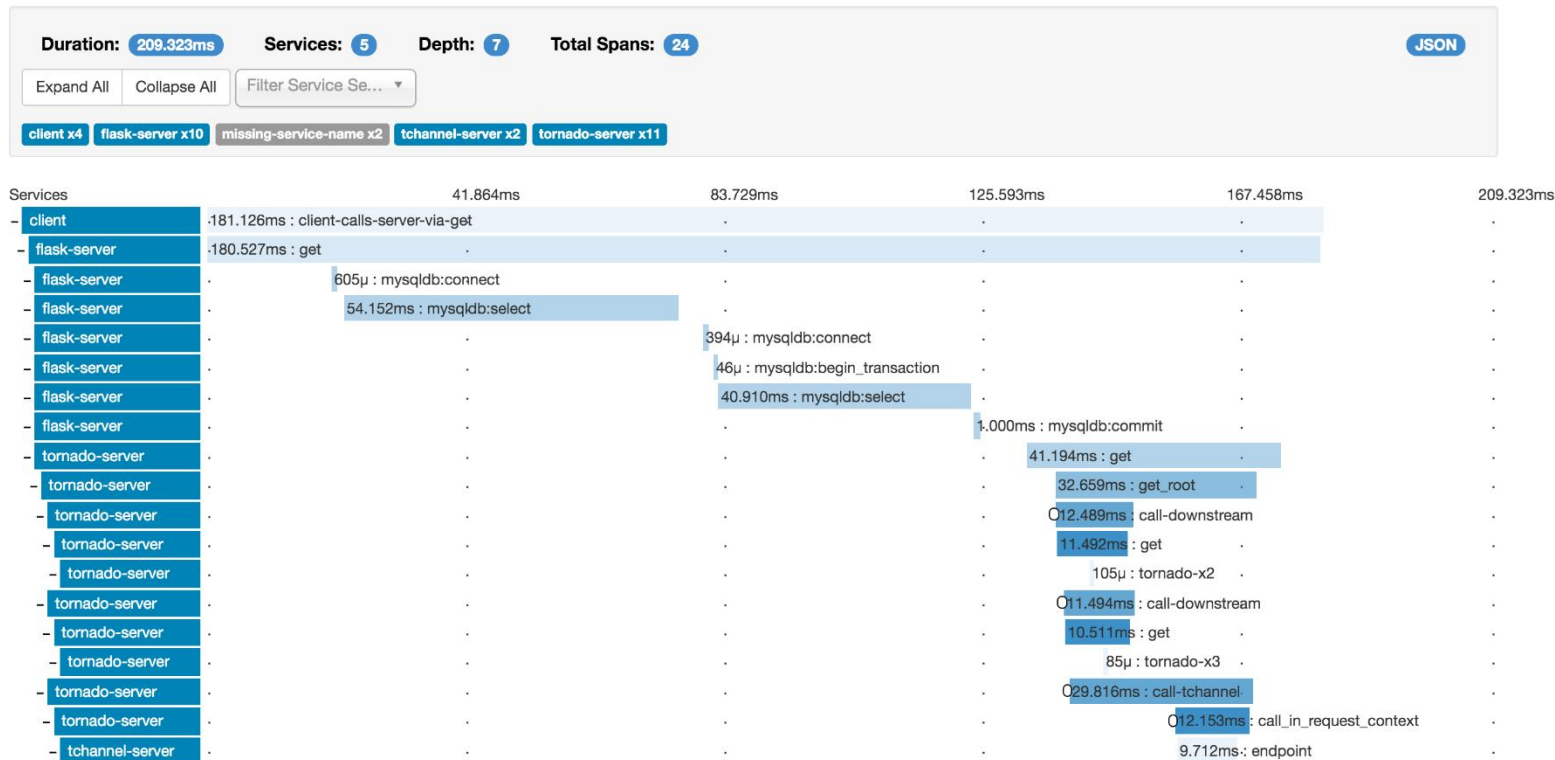
```
...
try (Scope scope = tracer.spanBuilder("main").startScopedSpan()) {
    System.out.println("About to do some busy work...");
    for (int i = 0; i < 10; i++) {
        doWork(i);
    }
}
```

For each individual span

```
...
public void doWork(int i) {
    // Starts another span, which will be a child span if another span is already active
    try (Scope scope = tracer.spanBuilder("main").startScopedSpan()) {
        // work
    }
}
```

Source: <https://opencensus.io/quickstart/java/tracing/>

ZIPKIN (zipkin.io)



Source: <https://zipkin.io/public/img/web-screenshot.png>

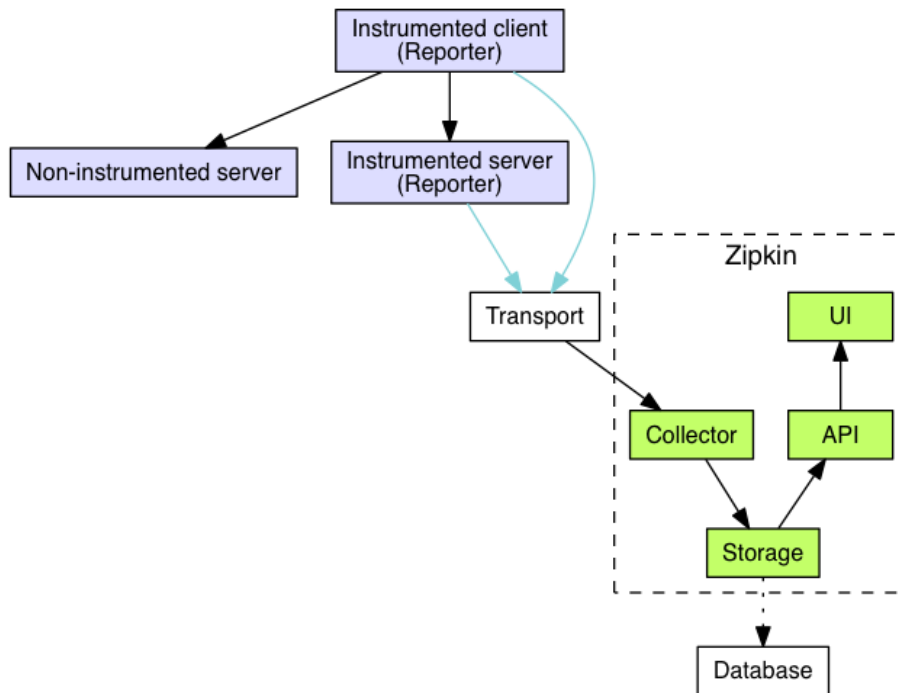
ZIPKIN (zipkin.io)

Supported Languages:

C#, Go, Java, JavaScript, Ruby,
Scala, PHP

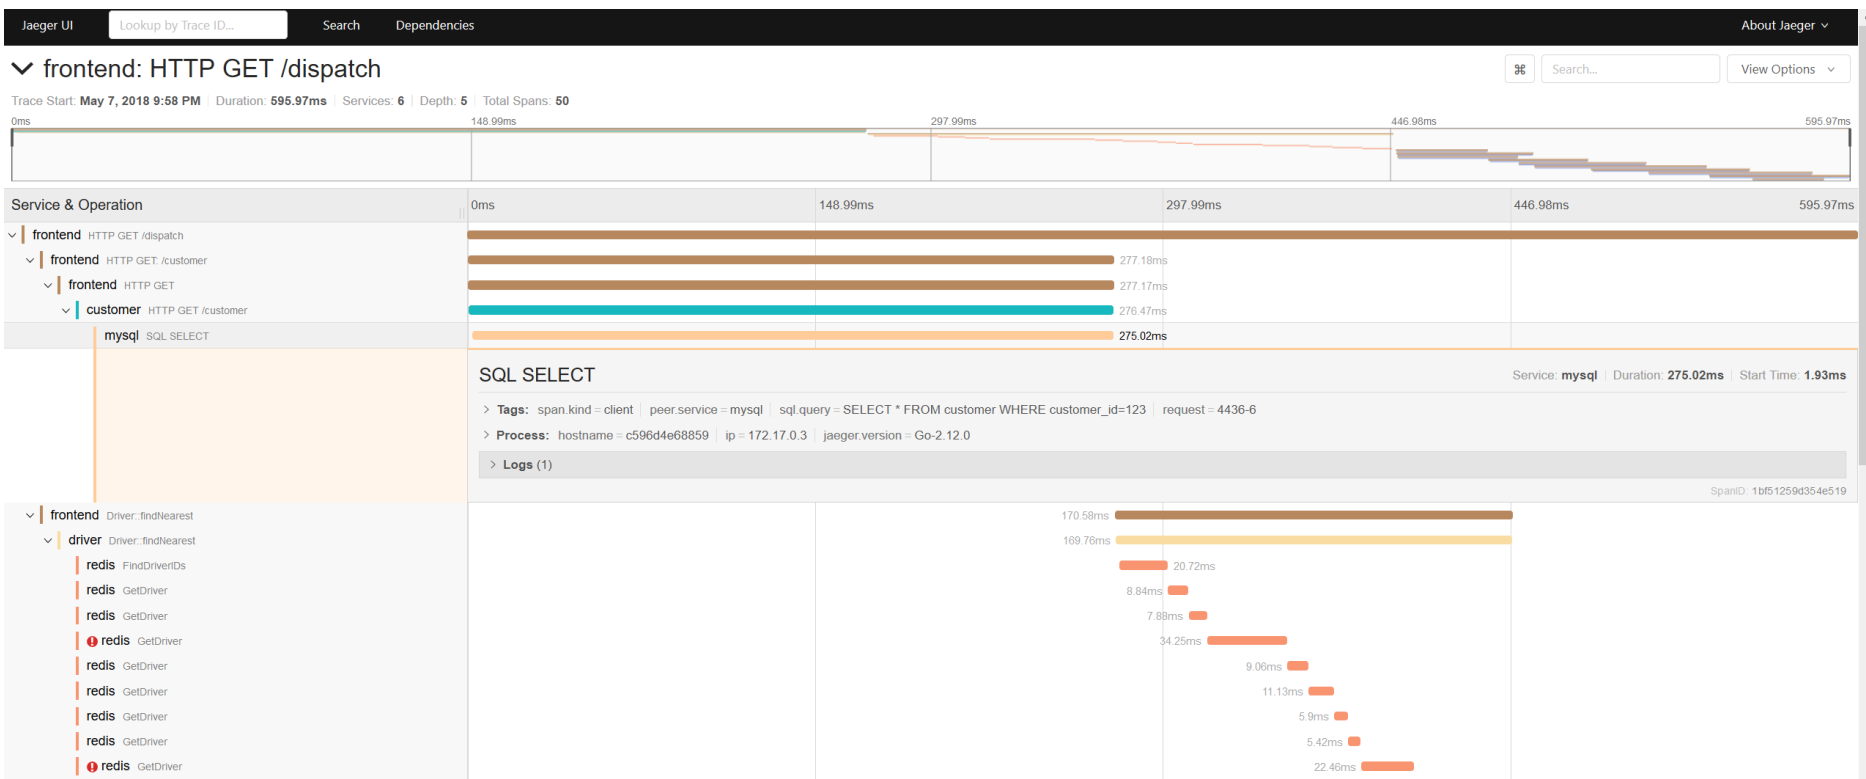
Supported Languages (Community Contributions):

C, C++, Elixir, Python, Scala, PHP



Source: <https://zipkin.io/pages/architecture.html>

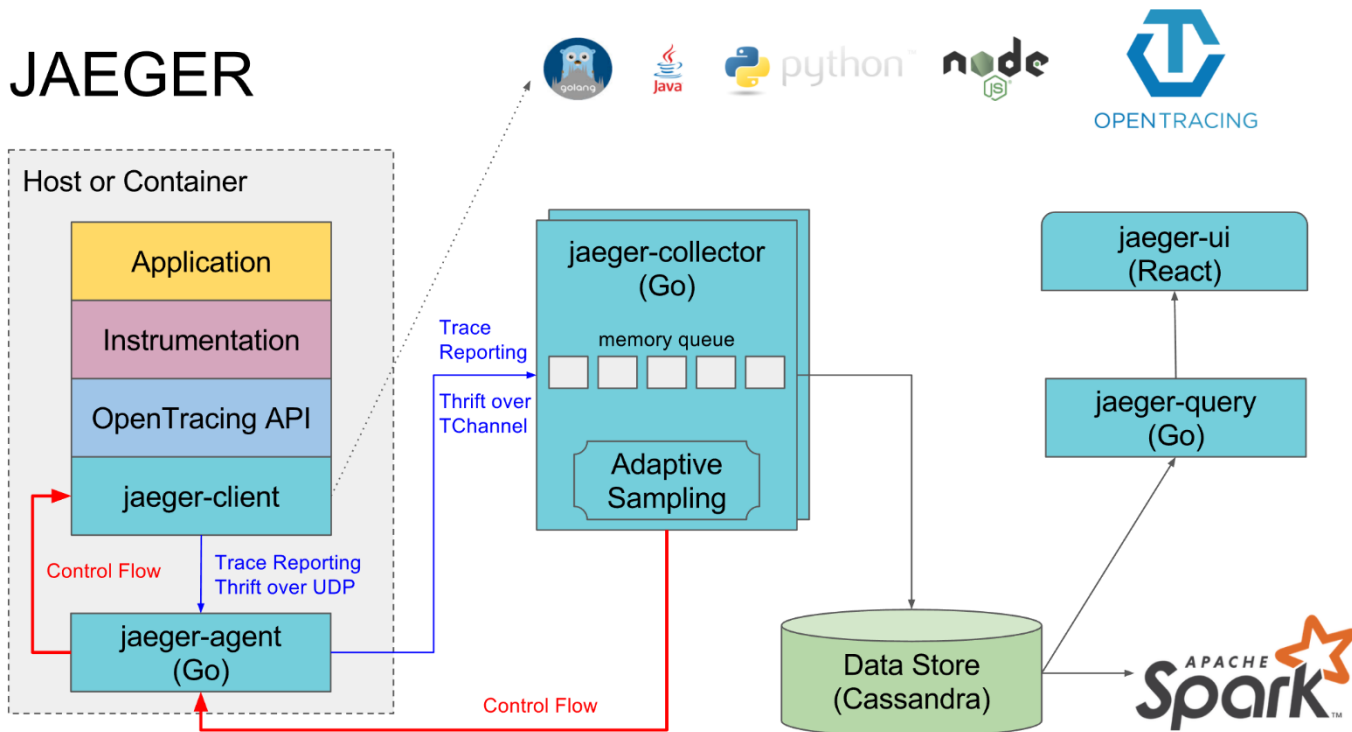
Jaeger (jaegertracing.io)



Jaeger (jaegertracing.io)

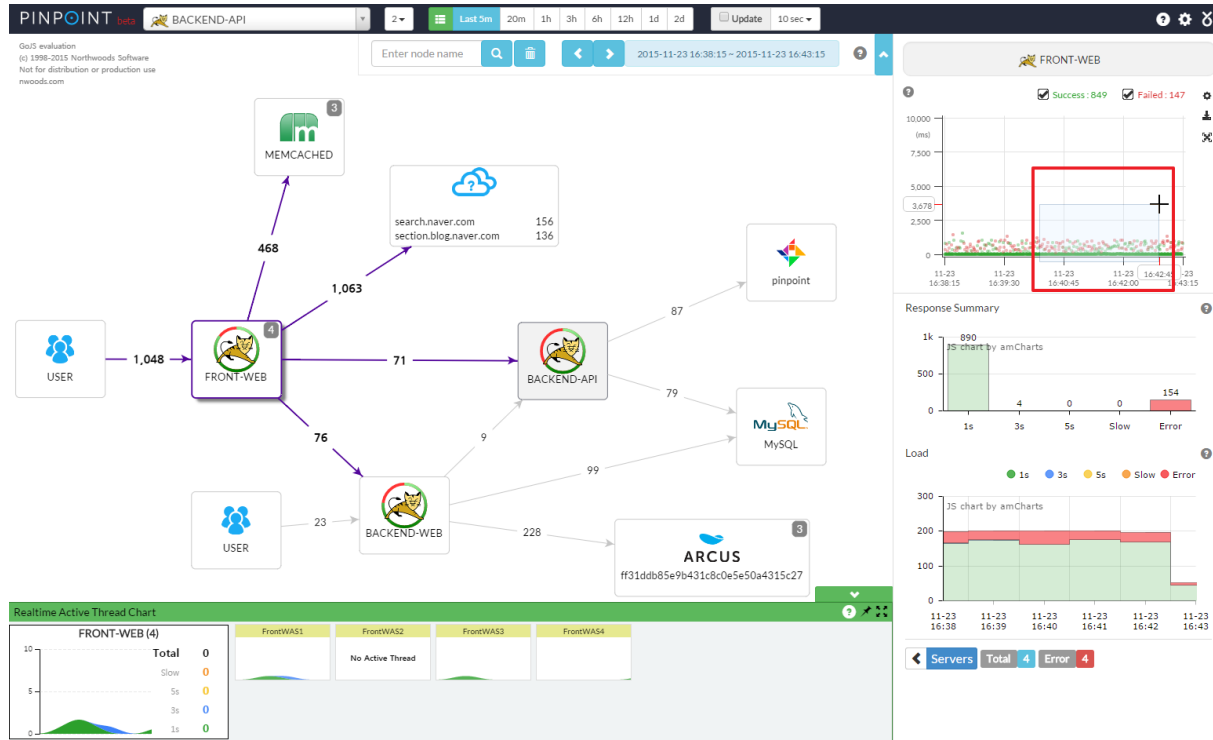
OPENTRACING

JAEGER



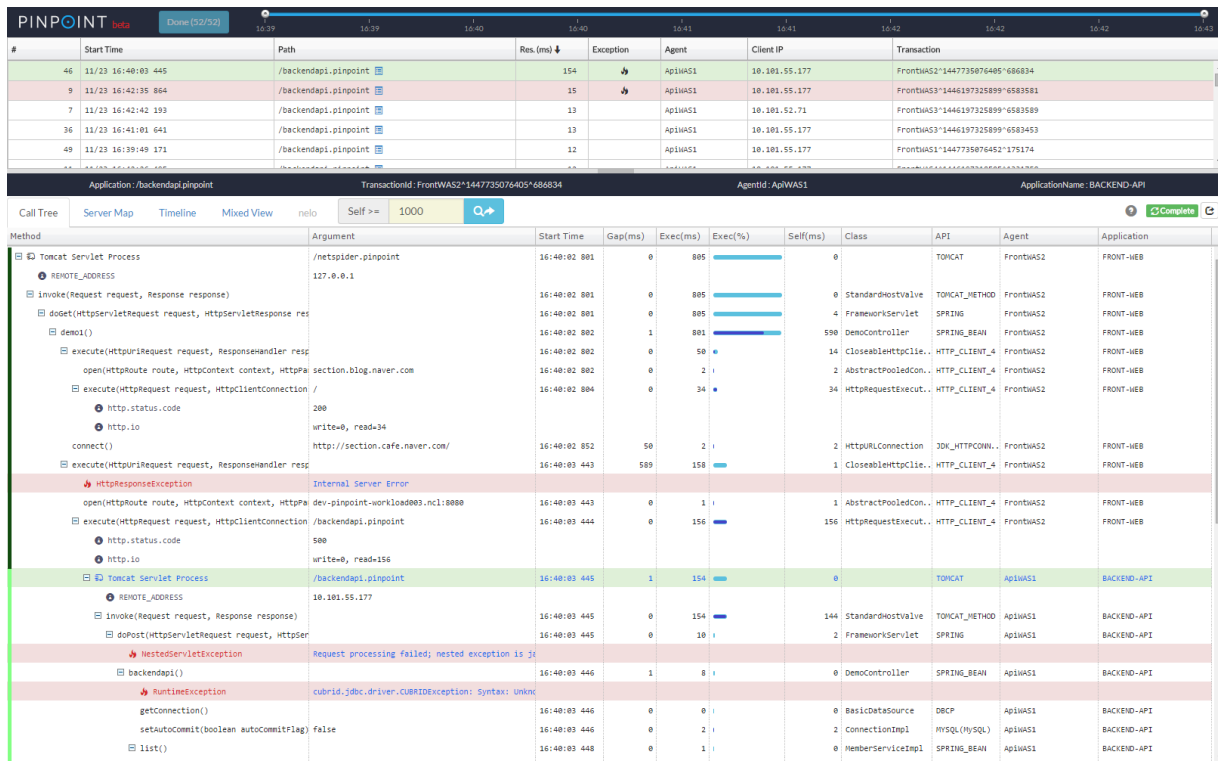
Source: <https://www.jaegertracing.io/docs/architecture/>

PINPOINT (http://naver.github.io/pinpoint/)



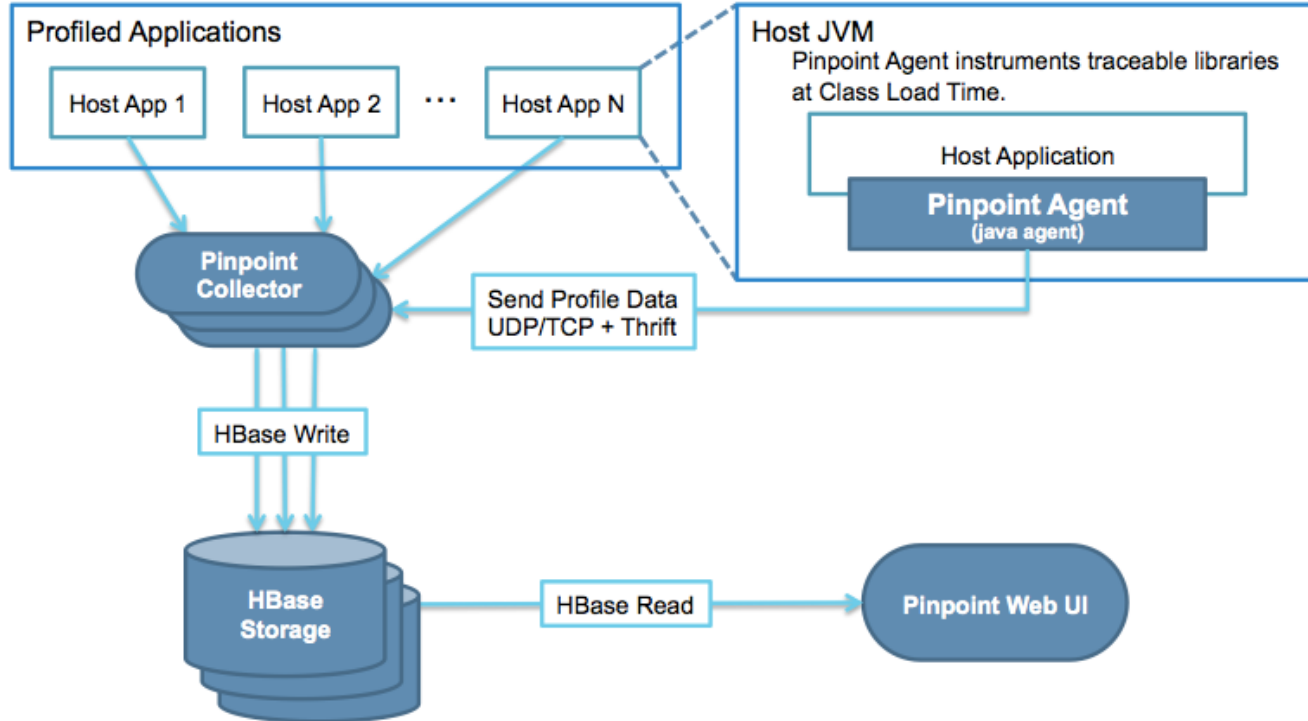
Source: <http://naver.github.io/pinpoint/overview.html>

PINPOINT (http://naver.github.io/pinpoint/)



Source: <http://naver.github.io/pinpoint/overview.html>

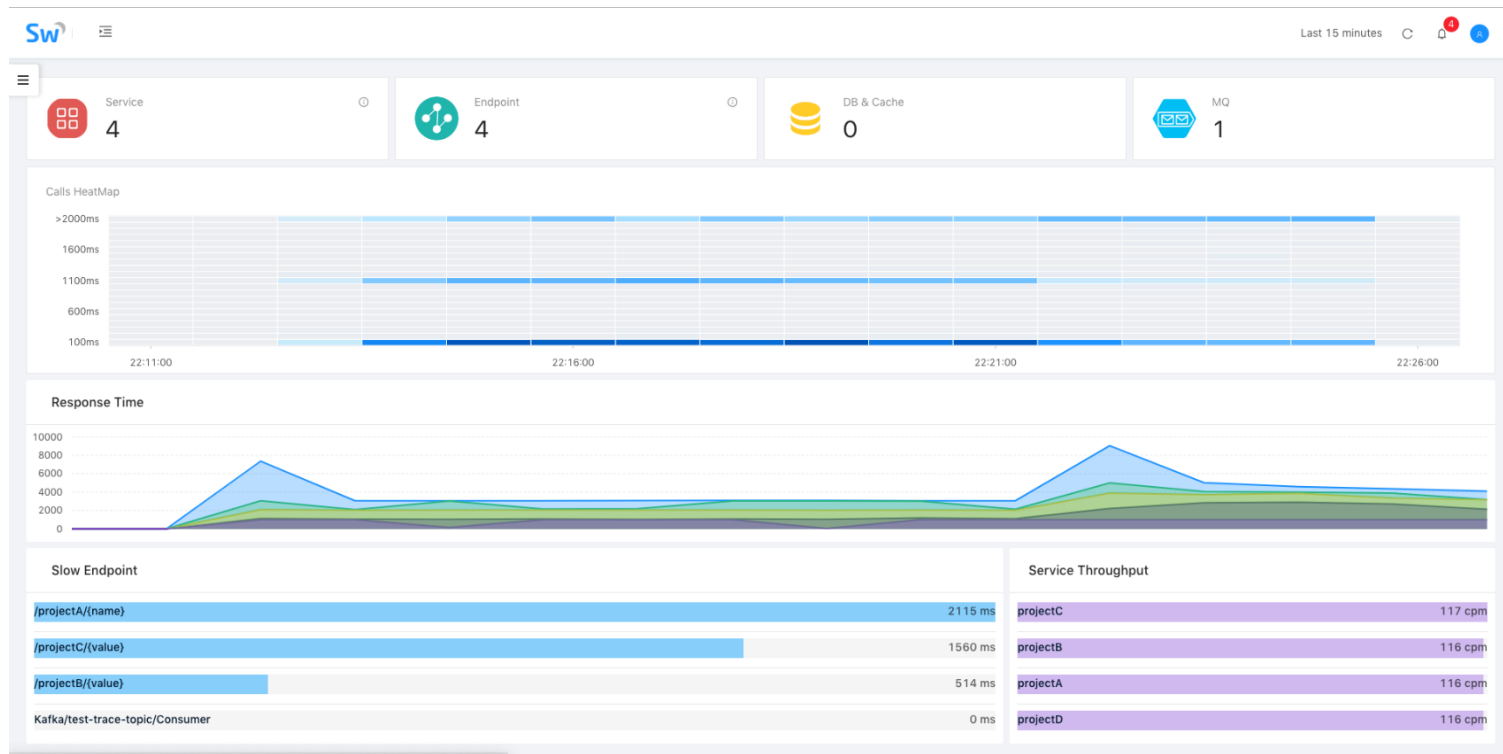
PINPOINT (<http://naver.github.io/pinpoint/>)



Source: <http://naver.github.io/pinpoint/overview.html>

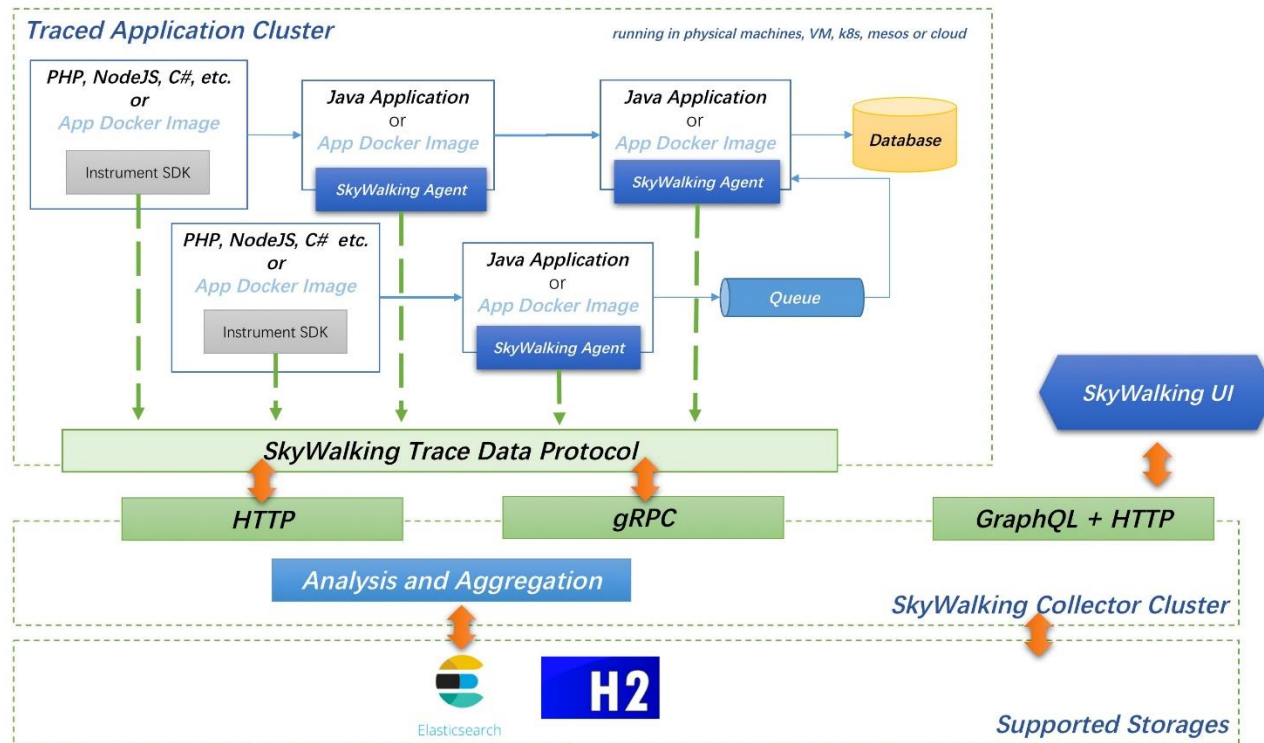
Apache Skywalking (skywalking.apache.org)

OPENTRACING



Source: <https://github.com/apache/incubator-skywalking/blob/master/docs/Screenshots.md#agent>

Apache Skywalking (skywalking.apache.org)



Source: <https://github.com/apache/incubator-skywalking>

Apache Skywalking (skywalking.apache.org)

OPENTRACING

Agent for Java, Instrumentation SDK for PHP, C#, NodeJS

HTTP Server

[Tomcat](#) 7
[Tomcat](#) 8
[Tomcat](#) 9
[Spring Boot](#) Web 4.x
Spring MVC 3.x, 4.x with servlet 3.x
[Nutz Web Framework](#) 1.x
[Struts2 MVC](#) 2.3.x -> 2.5.x
[Resin](#) 3 (Optional!)
[Resin](#) 4 (Optional!)
[Jetty Server](#) 9

HTTP Client

[Feign](#) 9.x
[Netflix Spring Cloud Feign](#) 1.1.x, 1.2.x, 1.3.x
[Okhttp](#) 3.x
[Apache httpcomponent HttpClient](#) 4.2, 4.3
[Spring RestTemplate](#) 4.x
[Jetty Client](#) 9
[Apache httpcomponent AsyncClient](#) 4.x

JDBC

MySQL Driver 5.x, 6.x
Oracle Driver (Optional!)
H2 Driver 1.3.x -> 1.4.x
[Sharding-JDBC](#) 1.5.x
PostgreSQL Driver 8.x, 9.x, 42.x

RPC Frameworks

[Dubbo](#) 2.5.4 -> 2.6.0
[Dubbox](#) 2.8.4
[Motan](#) 0.2.x -> 1.1.0
[gRPC](#) 1.x
[Apache ServiceComb Java Chassis](#) 0.1 -> 0.5, 1.0.x

MQ

[RocketMQ](#) 4.x
[Kafka](#) 0.11.0.0 -> 1.0

NoSQL

Redis
[Jedis](#) 2.x
[MongoDB Java Driver](#) 2.13-2.14, 3.3+
Memcached Client
[Spymemcached](#) 2.x
[Xmemcached](#) 2.x

Service Discovery

[Netflix Eureka](#)

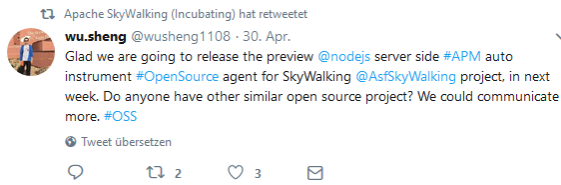
Spring Ecosystem

Spring Bean annotations (@Bean, @Service, @Component, @Repository) 3.x and 4.x (Optional!)
Spring Core Async SuccessCallback/FailureCallback/ListenableFutureCallback 4.x
[Hystrix: Latency and Fault Tolerance for Distributed Systems](#) 1.4.20 -> 1.5.12

Scheduler

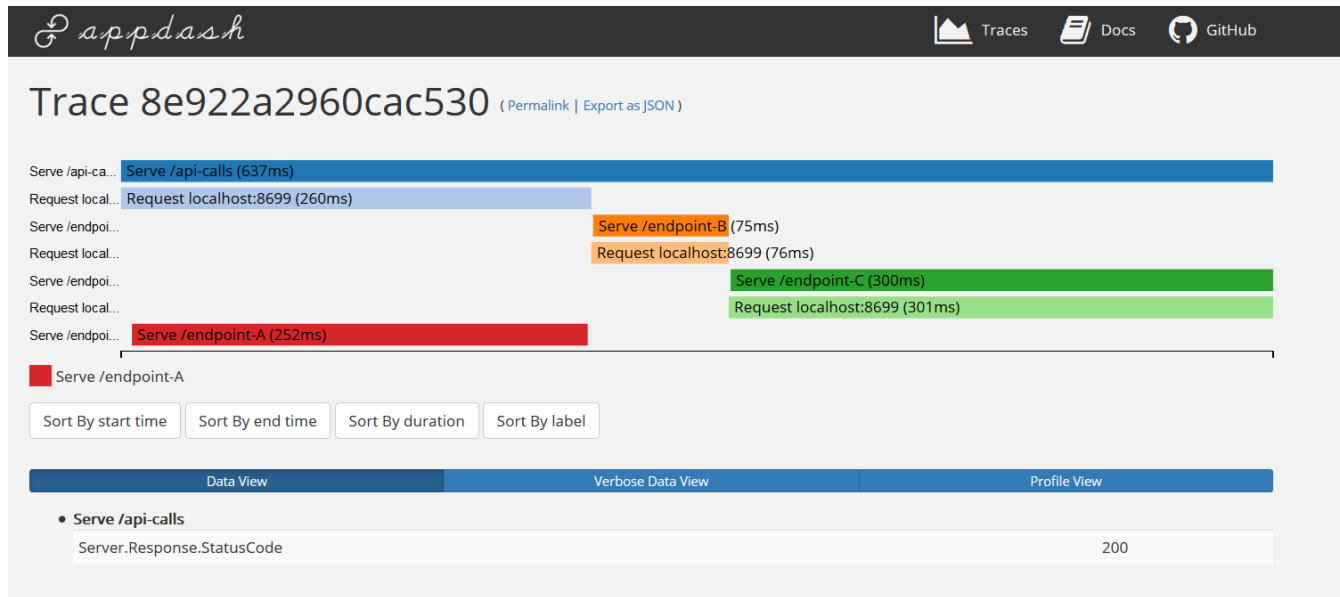
[Elastic Job](#) 2.x

OpenTracing community supported



AppDash (github.com/sourcegraph/appdash)

OPENTRACING



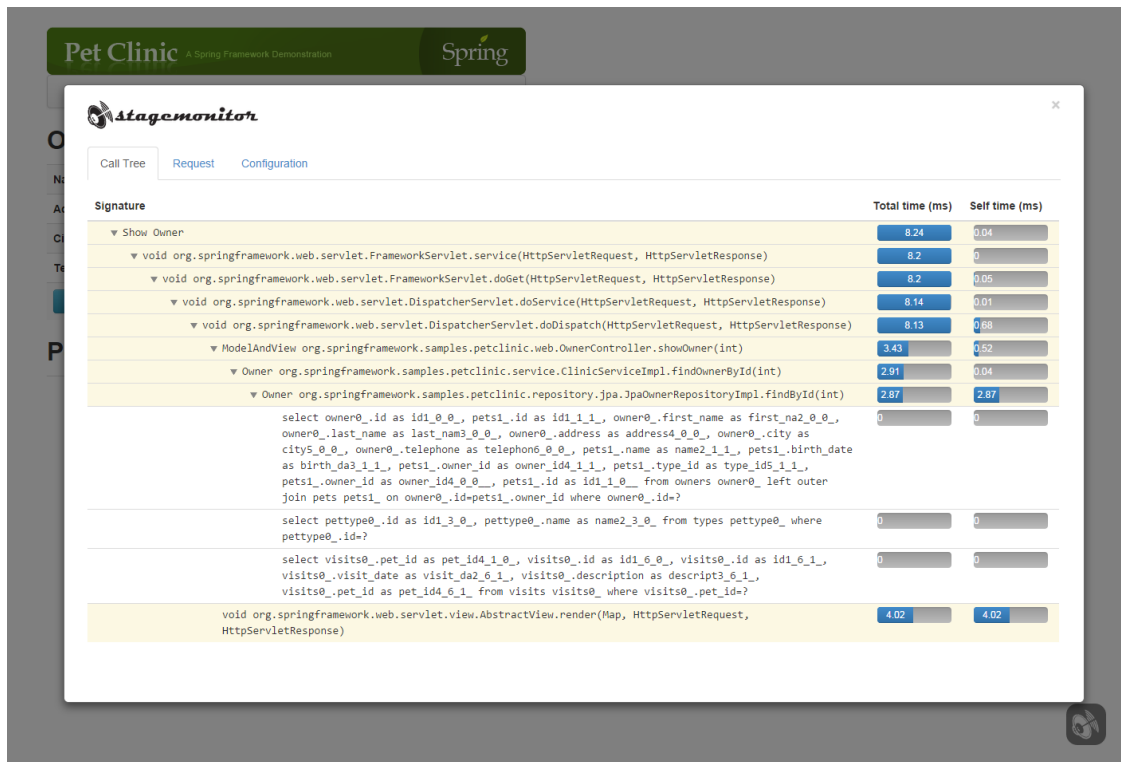
Supported Modules:

Go (<https://medium.com/opentracing/distributed-tracing-in-10-minutes-51b378ee40f1> ,

(Python - <https://github.com/sourcegraph/appdash/tree/master/python>),

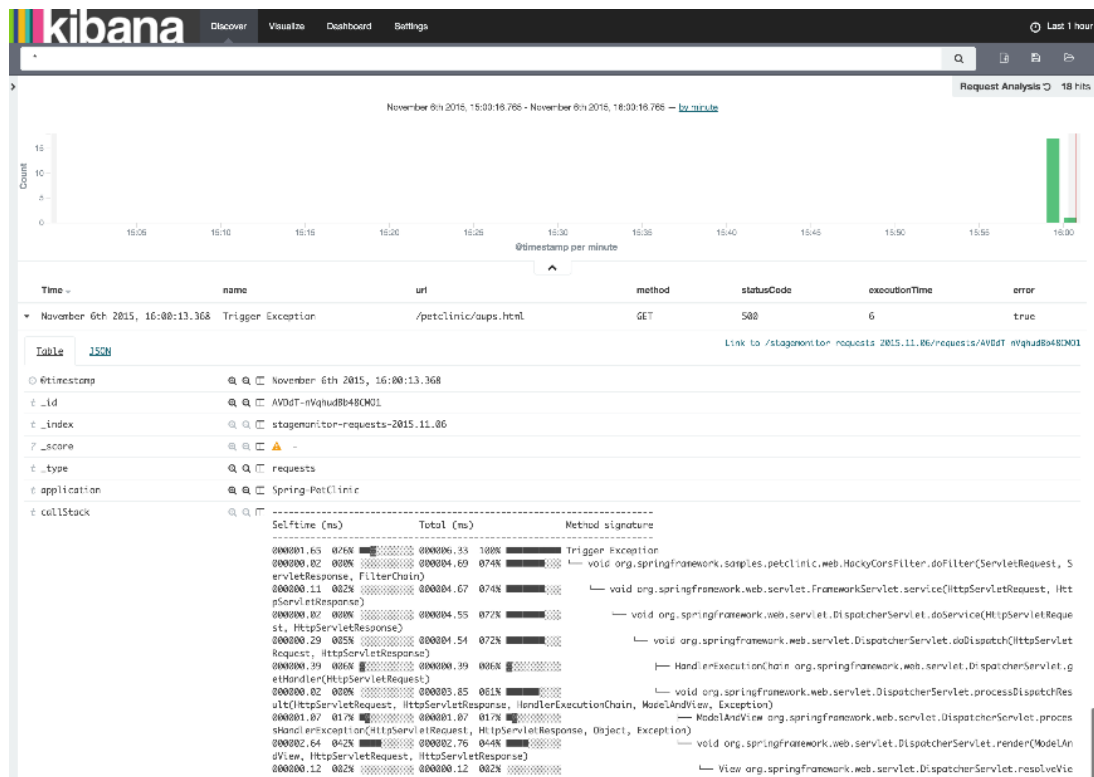
(Ruby - <https://github.com/bsm/appdash-rb>)

Stagemonitor (www.stagemonitor.org)



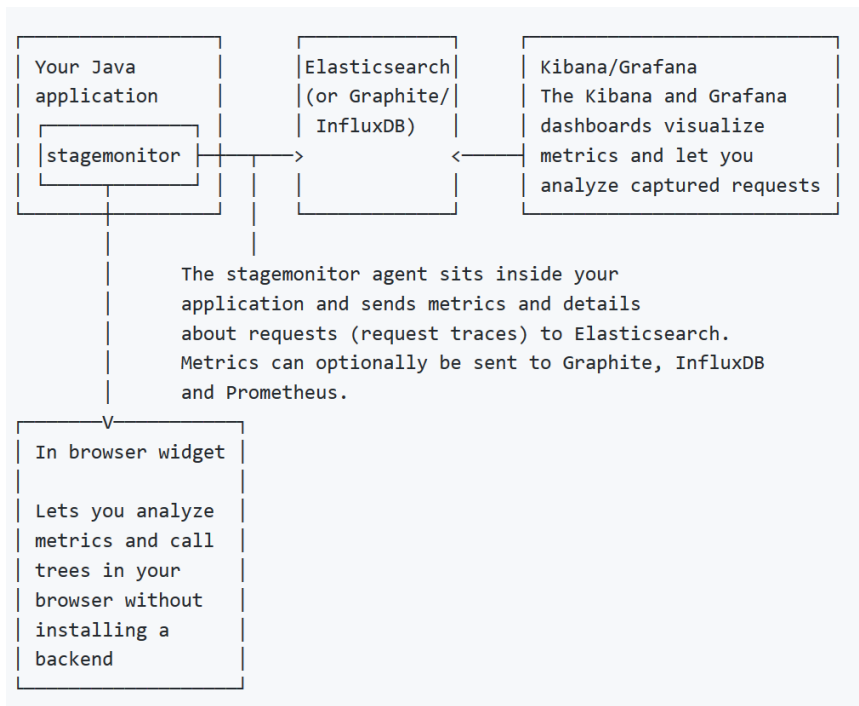
Source: <http://www.stagemonitor.org/de/#overview>

Stagemonitor (www.stagemonitor.org)



Source: <https://github.com/stagemonitor/stagemonitor/wiki/Request-Analysis-Dashboard>

Stagemonitor (www.stagemonitor.org)



Supported Modules:

Java (<https://github.com/stagemonitor/stagemonitor/wiki>)

InspectIT (inspectit.rocks)

The screenshot displays the InspectIT application interface. On the left, a sidebar shows the 'Local CMR' tree with 'inspectITDemo [n/a]' selected. The main window is titled 'Local CMR' and shows 'inspectITDemo [n/a]' with 'Invocation Sequences' expanded. A table lists invocation sequences with columns: Start Time, Method, Duration (ms), Child Count, URI, and Use case.

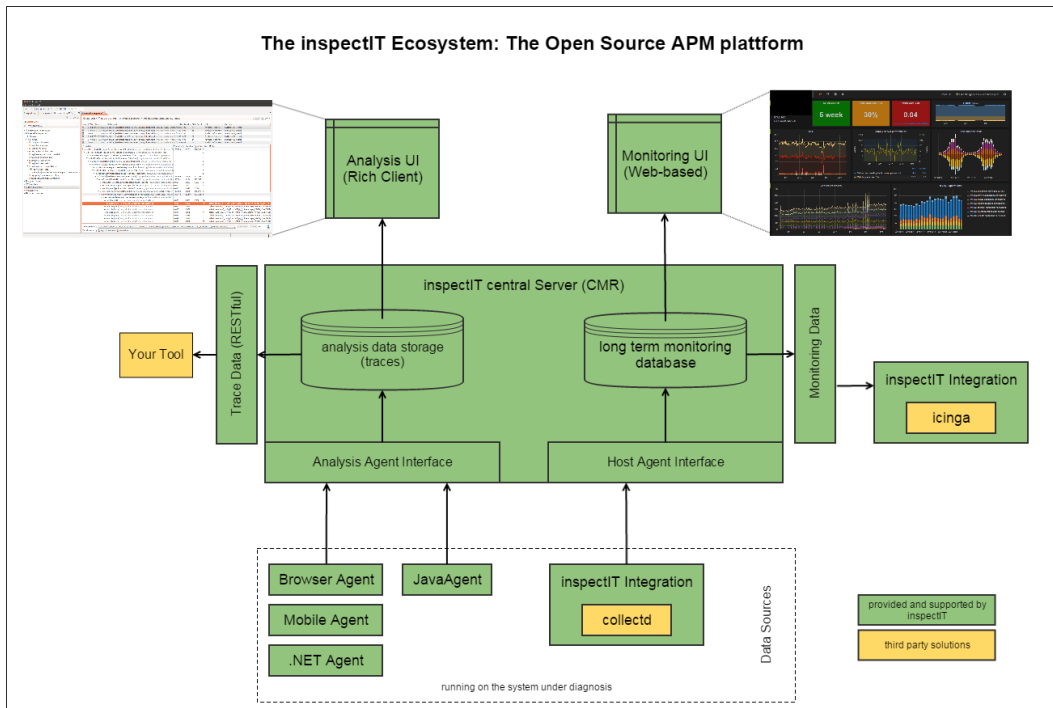
Start Time	Method	Duration (ms)	Child Count	URI	Use case
04.09.2012 17:20:04.43	doFilter(ServletRequest, ServletResponse, FilterChain) - org.jboss.web.tomcat.filters.Repl	437.868	84	/dvdstore/checkout	
04.09.2012 17:20:00.25	doFilter(ServletRequest, ServletResponse, FilterChain) - org.jboss.web.tomcat.filters.Repl	4172.216	122	/dvdstore/checkout	
04.09.2012 17:19:59.83	doFilter(ServletRequest, ServletResponse, FilterChain) - org.jboss.web.tomcat.filters.Repl	399.276	78	/dvdstore/browse	
04.09.2012 17:19:55.03	doFilter(ServletRequest, ServletResponse, FilterChain) - org.jboss.web.tomcat.filters.Repl	4777.570	177	/dvdstore/browse	Search
04.09.2012 17:19:52.32	doFilter(ServletRequest, ServletResponse, FilterChain) - org.jboss.web.tomcat.filters.Repl	2668.745	128	/dvdstore/browse	
04.09.2012 17:19:50.10	doFilter(ServletRequest, ServletResponse, FilterChain) - org.jboss.web.tomcat.filters.Repl	2212.059	40	/dvdstore/home	
04.09.2012 17:19:38.80	doFilter(ServletRequest, ServletResponse, FilterChain) - org.jboss.web.tomcat.filters.Repl	11216.400	25	/dvdstore/home	Home
04.09.2012 17:19:38.22	doFilter(ServletRequest, ServletResponse, FilterChain) - org.jboss.web.tomcat.filters.Repl	430.847	9	/dvdstore/	

Below the table, a detailed 'Method' call hierarchy is shown for the selected invocation. It includes columns: Method, Duration (ms), Exc. duration (s), Cpu Duration (s), Start Delta (ms), and SQL. The hierarchy shows the flow from the application's doFilter method through various filters and the servlet container to the database service.

Method	Duration (ms)	Exc. duration (s)	Cpu Duration (s)	Start Delta (ms)	SQL
doFilter(ServletRequest, ServletResponse, FilterChain) - org.jboss.web.tomcat.filters.ReplyHeaderFilter	4777.570	46.982	330.000	0	
doFilter(ServletRequest, ServletResponse, FilterChain) - org.jboss.seam.servlet.SeamFilter				0	
doFilter(ServletRequest, ServletResponse, FilterChain) - org.jboss.seam.web.HotDeployFilter				0	
doFilter(ServletRequest, ServletResponse, FilterChain) - org.jboss.seam.web.RedirectFilter				46	
doFilter(ServletRequest, ServletResponse, FilterChain) - org.jboss.seam.web.ExceptionFilter				46	
doFilter(ServletRequest, ServletResponse, FilterChain) - org.jboss.seam.web.MultipartFilter				46	
doFilter(ServletRequest, ServletResponse, FilterChain) - org.jboss.seam.web.IdentityFilter				46	
doFilter(ServletRequest, ServletResponse, FilterChain) - org.jboss.seam.web.LoggingFilter				46	
doFilter(ServletRequest, ServletResponse, FilterChain) - org.tuckey.web.filters.urlrew				47	
forward(ServletRequest, ServletResponse) - org.apache.catalina.core.ApplicationD	4730.588	0.230	290.000	47	
doForward(ServletRequest, ServletResponse) - org.apache.catalina.core.Applic	4730.358	2.096	290.000	47	
checkSameObjects(ServletRequest, ServletResponse) - org.apache.catalina.c	0.039	0.039	0.000	47	
wrapResponse(ApplicationDispatcher\$State) - org.apache.catalina.core.Appl	0.038	0.038	0.000	47	
wrapRequest(ApplicationDispatcher\$State) - org.apache.catalina.core.Appl	0.053	0.053	0.000	47	
processRequest(ServletRequest, ServletResponse, ApplicationDispatcher\$S	4728.131	0.122	290.000	47	
invoke(ServletRequest, ServletResponse, ApplicationDispatcher\$State) -	4728.009	399.389	290.000	47	
service(ServletRequest, ServletResponse) - javax.faces.webapp.FacesS				47	
unwrapRequest(ApplicationDispatcher\$State) - org.apache.catalina.co	0.107	0.107	0.000	4775	
unwrapResponse(ApplicationDispatcher\$State) - org.apache.catalina.c	0.038	0.038	0.000	4775	
recycleRequestWrapper(ApplicationDispatcher\$State) - org.apache.cal	0.038	0.038	0.000	4775	

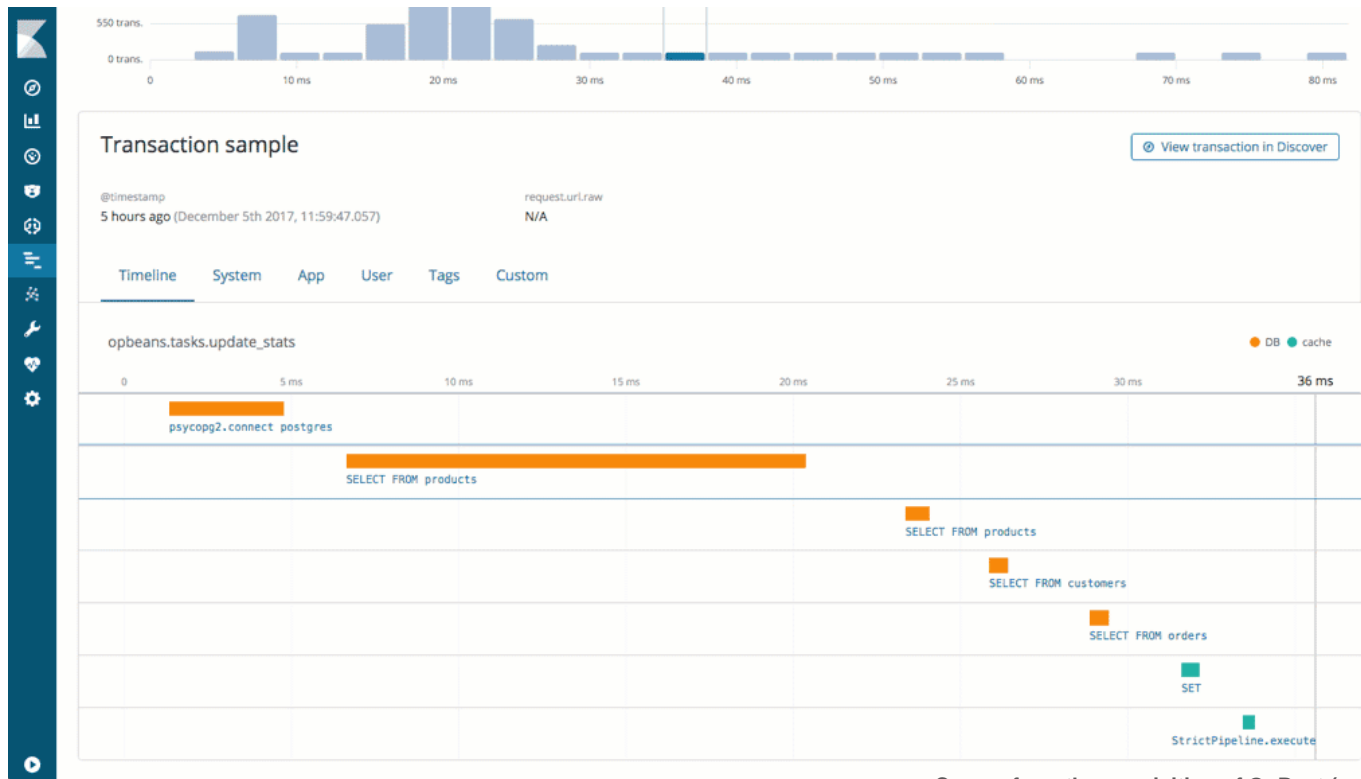
Source: <https://inspectit-performance.atlassian.net/wiki/spaces/DOC18/pages/93009319/Working+with+invocation+sequences>

InspectIT (inspectit.rocks)



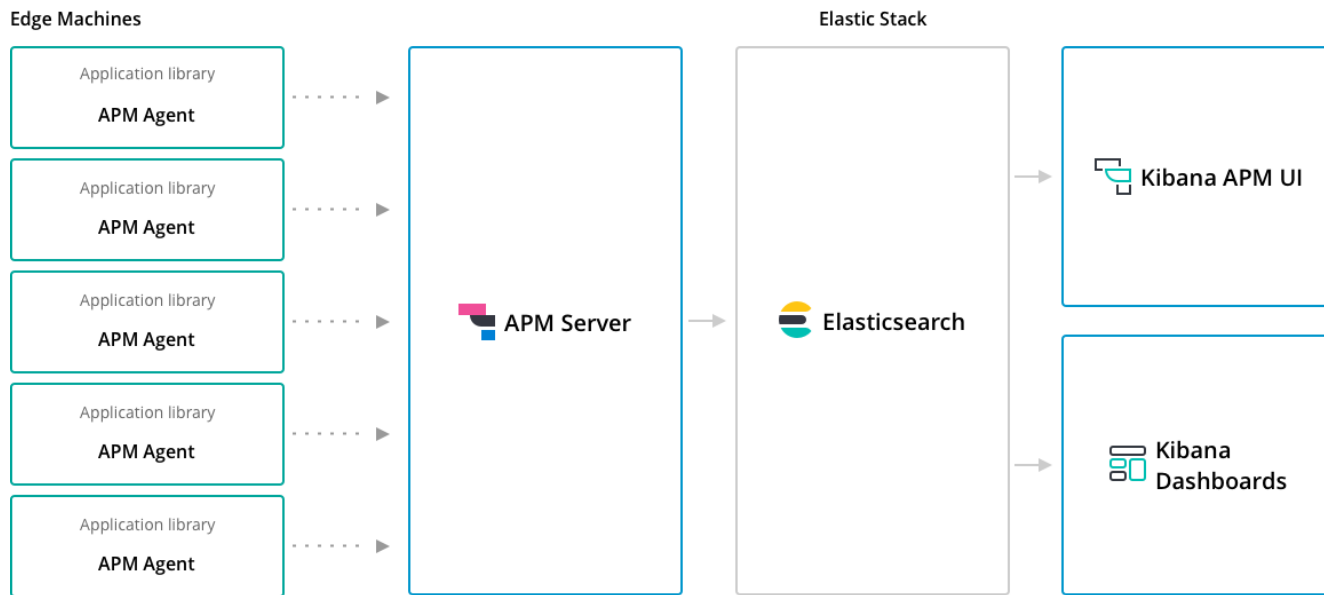
Supported Languages: Java, (.NET)

Elastic APM (www.elastic.co/solutions/apm)



Comes from the acquisition of OpBeat (part of Elastic Stack from 6.2):
<https://www.elastic.co/de/blog/elastic-apm-ga-released>

Elastic APM (www.elastic.co/solutions/apm)



Agents: Node.js, Python, Ruby, JavaScript, Go, Java, .NET
(<https://www.elastic.co/guide/en/apm/agent/index.html>)

Source: <https://www.elastic.co/guide/en/apm/get-started/current/overview.html>

What are reasons for a proprietary alternative?

- There is also cost associated with setting up and maintaining an open source APM solution (taken from <https://sematext.com/blog/performance-monitoring-comparison-build-vs-buy/>) :
 - **Build Your Own Monitoring System — Cost Scenario**
 - Hourly rate: 100 € (ballpark figure; could be much higher)
 - Installation: 2 hours (very optimistic)
 - Configuration: 8 hours (very optimistic)
 - Maintenance: 2 hours/month (optimistic)
 - Upgrading: 2 days (i.e., ~20 hours)/year (IF all goes well!)
 - # of servers to run this configuration: 3 (monitoring 10 total servers*)
 - Cost per server (hardware): 1,000 € each (i.e., 3,000 € total)
-
- Total Cost in Year 1: 6,200 €
 - Total Cost in Year 2: 3,200 € (not including any additional server purchases)
 - Total Cost in Year 3: 3,200 € (at least, though most likely higher)

What are reasons for a proprietary alternative?

- **Easier problem resolution:**
 - You do have someone to investigate and fix issues
 - Less risk in production as tools are (mostly) more thoroughly tested
- **Broader technology support:**
 - Developing agents is very time consuming and, thus, costly – the open source community cannot spend the same amount of manpower into this effort for each and every version of a technology (e.g., supporting Tomcat, 5,6,7,8, ...)
- **You can plan ahead:**
 - Vendors typically communicate the time until which a software version is supported and support the transition phase as well, this is not always the case for open source software

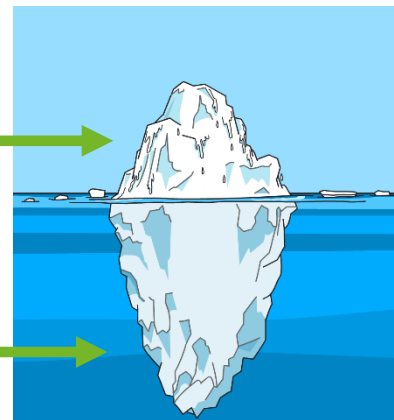
What are reasons for a proprietary alternative?

Remember: Code and Effort distribution of an APM Solution

UI + Server + Collectors



Agents



- Some things might change, as some open source projects (e.g., istio/Ingres/WildFly) are already supporting OpenTracing natively
- Furthermore, there are default implementations for Spring Boot or Thorntail (previously WildFly Swarm) to automatically capture traces that can be packaged in your application

Thank you!

Dr. Andreas Brunnert

brunnert@retit.de

Bernhard Lubomski

lubomski@retit.de



Resource Efficient Technologies & IT Systems